

# **Walmart Sales Trends Predictions and Analysis via Machine Learning**

**Utilizing Advanced Machine Learning Techniques for  
Sales Insights**

Umar Memon and Hamza Khan



# Introduction and Objectives

## Background and Motivation:

- For our project we wanted to focus on analyzing data and make predictions on it. We chose this idea because after researching about machine learning projects we wanted to do a project that would be helpful in our career and we could learn from.

## Objective:

- Our main object was to analyze and make sales or demand trend predictions on retail stores.
- After researching we understood why predictions, forecasting, analyzing was important, it's for business strategies and most importantly looking at trends of product demands from a retail store.

# Data Collection and Preprocessing

## Dataset Description:

**Source:** Found one dataset of Walmart sales, States, Years, and other main features from data world website. We did see other Walmart dataset however this one has more features, more rows, and has data for US only. This dataset was about 300MB in size so it had a lot of data.

**Key Features:** Sales, unit prices, quantity order, discount, profit, and other relevant features.

## Preprocessing Steps:

- Handled missing value by doing imputation techniques. Removed outliers using the IQR (interquartile range) method from statistic and probability. Used stratified random sampling to only use 1000 samples from the dataset.
- Feature Engineering, we created new features such as estimated sale using other important feature like profit, unit price, discount, sales, quantity order, any features that gave more insight in the demand of products using linear combination method. Also created log sales, log prices, and more.

# Data Collection and Preprocessing

```
2 print(dataset.head())
```

	city	customer_age	customer_name	customer_segment	discount	\
128105	194.0	76.0	101.0	1.0	0.08	
225930	77.0	65.0	475.0	1.0	0.22	
355800	422.0	80.0	309.0	2.0	0.18	
47215	300.0	45.0	325.0	0.0	0.22	
161123	95.0	85.0	125.0	2.0	0.05	

	order_date	order_id	order_priority	order_quantity	\
128105	2022-09-27	194.0	0.0	20.0	
225930	2022-03-19	988.0	2.0	32.0	
355800	2021-08-28	392.0	4.0	3.0	
47215	2021-03-23	376.0	1.0	4.0	
161123	2022-04-30	686.0	2.0	16.0	

	product_base_margin	...	profit_margin	order_date_dayofweek	\
128105	0.49	...	58.735697	1	
225930	0.40	...	21.884385	5	
355800	0.57	...	49.777919	5	
47215	0.38	...	2.756167	1	
161123	0.59	...	0.522490	5	

	order_date_month	order_date_year	order_date_dayssincestart	\
128105	9	2022	1364	
225930	3	2022	1172	
355800	8	2021	969	
47215	3	2021	811	
161123	4	2022	1214	

	ship_date_dayofweek	ship_date_month	ship_date_year	\
128105	1	10	2022	
225930	2	3	2022	
355800	2	9	2021	
47215	2	3	2021	
161123	3	5	2022	

	ship_date_dayssincestart	estimated_sales
128105	1368	2355.631031
225930	1173	2178.492221
355800	970	1240.908999
47215	816	19.512609
161123	1216	150.862391

[5 rows x 38 columns]

First few rows of the dataset:

	city	customer_age	customer_name	customer_segment	discount	\
128105	Eulless	76	Christine Abelman	Corporate	0.08	
225930	Burlington	65	Sonia Cooley	Corporate	0.22	
355800	Moorhead	80	Kristen Hastings	Home Office	0.18	
47215	Hurst	45	Logan Currie	Consumer	0.22	
161123	Chattanooga	85	Cyma Kinney	Home Office	0.05	

	order_date	order_id	order_priority	\
128105	2022-09-27	34a99182-8b7a-4a61-a06d-5a61eb926f7d	Critical	
225930	2022-03-19	fc94e58e-62e2-4d51-bbcf-b29469fb7e1b	Low	
355800	2021-08-28	6491fc97-f7b8-4e5f-9104-1a499d586dbf	Not Specified	
47215	2021-03-23	5f5daae8-77bc-47c0-a7da-bd4b4be7af8f	High	
161123	2022-04-30	b2043249-8d7b-4aaa-836f-1c688a50091d	Low	

	order_quantity	product_base_margin	...	product_sub_category	\
128105	20	0.49	...	Office Furnishings	
225930	32	0.4	...	Paper	
355800	3	0.57	...	Appliances	
47215	4	0.38	...	Paper	
161123	16	0.59	...	Telephones and Communication	

	profit	region	sales	ship_date	ship_mode	\
128105	23376.80793	Central	398	2022-10-04	Express Air	
225930	21541.2385	East	984.32	2022-03-23	Express Air	
355800	12116.94153	Central	243.42	2021-09-01	Express Air	
47215	132.95754	Central	48.24	2021-03-31	Delivery Truck	
161123	1052.8381	South	2015.04	2022-05-05	Regular Air	

	shipping_cost	state	unit_price	zip_code
128105	17.759847984806232	Texas	19.98	76039.0
225930	659.3479188838185	Vermont	30.98	5401.0
355800	206.306215156242	Minnesota	81.32	56560.0
47215	19.83636696486954	Texas	12.28	76053.0
161123	710.1907635117352	Tennessee	125.99	37421.0

[5 rows x 23 columns]

# Data Collection and Preprocessing

```
44
45 split = StratifiedShuffleSplit(n_splits=1, test_size=n_samples, random_state=42)
46 for _, sample_index in split.split(filtered_df, filtered_df[stratify_column]):
47     stratified_sample = filtered_df.iloc[sample_index]
48
```

```
# Feature Engineering for Date Columns
for col in ['order_date', 'ship_date']:
    if col in dataset.columns:
        dataset[col + '_dayofweek'] = dataset[col].dt.dayofweek
        dataset[col + '_month'] = dataset[col].dt.month
        dataset[col + '_year'] = dataset[col].dt.year
        dataset[col + '_dayssincestart'] = (dataset[col] - dataset[col].min()).dt.days
```

```
105 # Calculate estimated sales based on the features
106 intercept = 0.1 # Placeholder value
107 coef_quantity = 0.4 # Placeholder value
108 coef_price = 0.3 # Placeholder value
109 coef_discount = 0.2 # Placeholder value
110 coef_profit = 0.1
111 coef_log_sales = 0.1
112 coef_discount_per_quantity = 0.05
113 coef_profit_margin = 0.05
114 coef_log_price = 0.1
115
116 dataset = dataset[dataset['order_date_year'] != 2023]
117
118 dataset['estimated_sales'] = abs(
119     intercept
120     + coef_quantity * dataset['order_quantity']
121     + coef_price * dataset['unit_price']
122     + coef_discount * dataset['discount']
123     + coef_profit * dataset['profit']
124     + coef_log_sales * dataset['log_sales']
125     + coef_discount_per_quantity * dataset['discount_per_quantity']
126     + coef_profit_margin * dataset['profit_margin']
127     + coef_log_price * dataset['log_price']
128 )
129
```

```
65 # Handle missing values
66 logging.info("Handling missing values...") # Log message for handling missing values
67 imputer = SimpleImputer(strategy='mean') # Initialize imputer to fill missing values with mean
68 numerical_columns = dataset.select_dtypes(include=['number']).columns # Select only numerical columns
69 dataset[numerical_columns] = imputer.fit_transform(dataset[numerical_columns]) # Apply imputer
70
71 # # Normalize/Standardize data
72 # logging.info("Normalizing/Standardizing data...")
73 # numerical_features = ['customer_age', 'discount', 'order_quantity', 'product_base_margin',
74 #                        'unit_price']
75 # scaler = StandardScaler()
76 # dataset[numerical_features] = scaler.fit_transform(dataset[numerical_features])
77
78 # Feature engineering: Create new features
79 logging.info("Creating new features...") # Log message for feature engineering
80 dataset['price_adjusted_by_quantity'] = dataset['unit_price'] * dataset['order_quantity']
81 dataset['log_price'] = np.log1p(dataset['unit_price'])
82 dataset['inverse_price'] = 1 / (dataset['unit_price'] + 1e-5)
83 dataset['log_sales'] = np.log1p(dataset['sales'])
84 dataset['discount_per_quantity'] = dataset['discount'] / (dataset['order_quantity'] + 1e-5)
85 dataset['profit_margin'] = dataset['profit'] / (dataset['sales'] + 1e-5)
86
```

```
2 # Function to check for negative values
3 def check_for_negative_values(dataset, stage):
4     logging.info(f"Checking for negative values after {stage}...") # Log message for checking negative values
5     numerical_features = dataset.select_dtypes(include=[np.number]).columns # Get numerical columns
6     for feature in numerical_features: # Iterate through numerical columns
7         if (dataset[feature] < 0).any(): # Check if there are negative values
8             dataset[feature] = abs(dataset[feature]) # Convert negative values to positive
9             logging.warning(f"Negative values found and converted to positive in {feature} after {stage}")
10
```

# (EDA) Exploratory Data Analysis

## Data Exploration:

**Libraries:** Imported various libraries, like pandas, numpy, seaborn, matplotlib, linear regression models, time series, and others for forecasting and visualizations.

**Data Cleaning:** Removed duplicates, corrected data types, and dropped rows/columns with missing data.

**Data Transformation:** Performed data transformation converting date column to datetime format for time series analysis.

**Descriptive Statistics:** Calculated statistics for numerical column, providing insight in dispersion, central tendency, and shape of the data distribution.

**Correlation Analysis and Distribution Visualization:** computed correlation matrix and visualized it using heatmap to see relationships between different features. Visualization distributions and relationships of numerical variables using box plots, bar graphs, scatter plots, and histograms.

# Exploratory Data Analysis

```
1 from sklearn.preprocessing import PolynomialFeatures
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 import logging
5 import numpy as np
6 from sklearn.linear_model import LinearRegression
7 from sklearn.svm import SVC
8 from sklearn.preprocessing import StandardScaler
9 import pandas as pd
10
```

```
1 from sklearn.linear_model import Lasso, LinearRegression
2 from sklearn.pipeline import Pipeline
3 from sklearn.preprocessing import PolynomialFeatures, StandardScaler, RobustScaler, MinMaxScaler, MaxAbsScaler
4 from sklearn.model_selection import RandomizedSearchCV, train_test_split, cross_val_score
5 from sklearn.metrics import mean_squared_error, r2_score
6 import matplotlib.pyplot as plt
7 import logging
8 import numpy as np
```

```
1 # Libraries
2 import pandas as pd
3 import numpy as np
4 from sklearn.model_selection import StratifiedShuffleSplit
5 import logging
6 import pandas as pd
7 import numpy as np
8 from google.colab import files, drive
9 from sklearn.preprocessing import StandardScaler, LabelEncoder
0 from sklearn.impute import SimpleImputer
1 from sklearn.linear_model import LinearRegression, Lasso
2 from sklearn.model_selection import train_test_split, RandomizedSearchCV
3 from sklearn.pipeline import Pipeline
4 from sklearn.metrics import mean_squared_error, r2_score
5 import logging
6 import matplotlib.pyplot as plt
7 import seaborn as sns
```

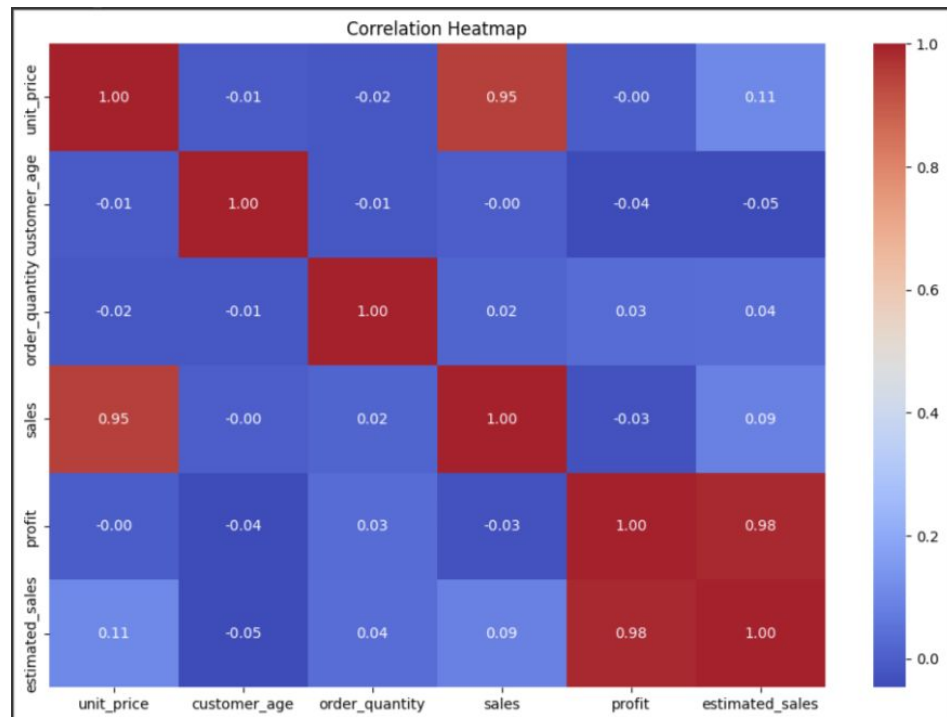

```
from statsmodels.tsa.statespace.sarimax import SARIMAX
from sklearn.model_selection import train_test_split
```

```
4 from prophet import Prophet
5 from prophet.plot import add_changepoints_t
```






# Exploratory Data Analysis (Demo)





# Exploratory Data Analysis (Demo)



# Exploratory Data Analysis

```
1 from sklearn.preprocessing import PolynomialFeatures, LabelEncoder
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 import logging
5 import numpy as np
6 import pandas as pd
7
```

```
8 # Function to calculate correlation
9 def calculate_correlation(data):
10     correlation_price_sales = data[['unit_price', 'estimated_sales']].corr()
11     logging.info(f"Correlation between Unit Price and Estimated Sales:\n{correlation_price_sales}")
12
13     correlation_sales_estimated = data[['sales', 'estimated_sales']].corr()
14     logging.info(f"Correlation between Sales and Estimated Sales:\n{correlation_sales_estimated}")
15
16     return correlation_price_sales, correlation_sales_estimated
17
18 # Function to visualize the relationship between features and estimated sales
19 def visualize_relationship(data):
20     plt.figure(figsize=(10, 6))
21     plot_df = pd.DataFrame({
22         'unit_price': data['unit_price'].values.ravel(),
23         'estimated_sales': data['estimated_sales'].values.ravel()
24     })
25     sns.scatterplot(x='unit_price', y='estimated_sales', data=plot_df)
26     plt.title('Relationship between Unit Price and Estimated Sales')
27     plt.xlabel('Unit Price')
28     plt.ylabel('Estimated Sales')
29     plt.show()
30
```

```
56 # Function to visualize the relationship between features and estimated sales
57 def visualize_relationship(data):
58
59     features = ['unit_price', 'order_quantity', 'discount', 'profit', 'sales']
60     for feature in features:
61         if feature in data.columns and 'estimated_sales' in data.columns: # Check if columns exist
62             plt.figure(figsize=(10, 6))
63             plot_df = pd.DataFrame({
64                 feature: data[feature].values.ravel(), # Ravel the array to ensure 1D
65                 'estimated_sales': data['estimated_sales'].values.ravel() # Ravel the array to ensure 1D
66             })
67             sns.scatterplot(x=feature, y='estimated_sales', data=plot_df) # Pass DataFrame to sns
68             plt.title(f'Relationship between {feature} and Estimated Sales')
69             plt.xlabel(feature)
70             plt.ylabel('Estimated Sales')
71             plt.show()
72         else:
73             logging.warning(f"Skipping visualization for {feature} due to missing column(s).")
74
75 # Ravel flattens a multi dimensional array
76
```

# Models

## Supervised Learning:

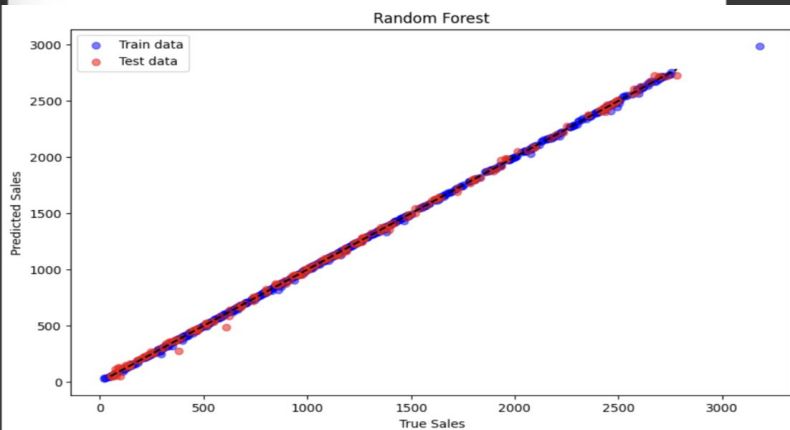
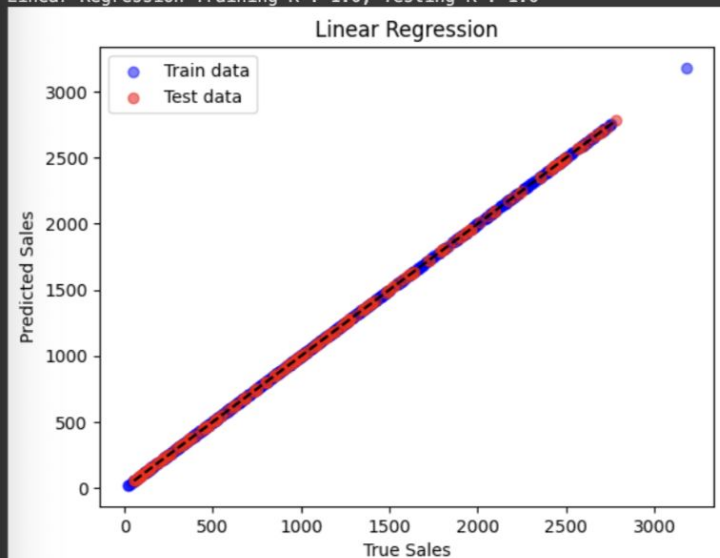
**Regression Models:** Linear Regression, Lasso Regression, Ridge Regression, Elastic, Random Forest, Gradient Boosting, Neural Network, and Ensemble model (combining all of these model for better performance).

**Training and Validation:** Performed training and testing split, hyperparameter tuning using Grid and Randomized Search to enhanced each models parameters. Cross-Validation techniques to ensure robust performance evaluations. Used pipeline, streamlining a machine learning process steps and param-grid defining set of hyperparameters tuning the models. Also did standardizing/Normalizing on models which had benefits.

**Performance of Regression Model and Forecasting:** Used  $R^2$  statistical measure to evaluate how well a regression model fit data. Also used MSE (Mean Squared Error) to evaluate the performance of predictive models. The closer the  $R^2$  is to 1 means the model is performing well, and if MSE has lower values then the model is performing well. Used Ensemble Model to forecast.

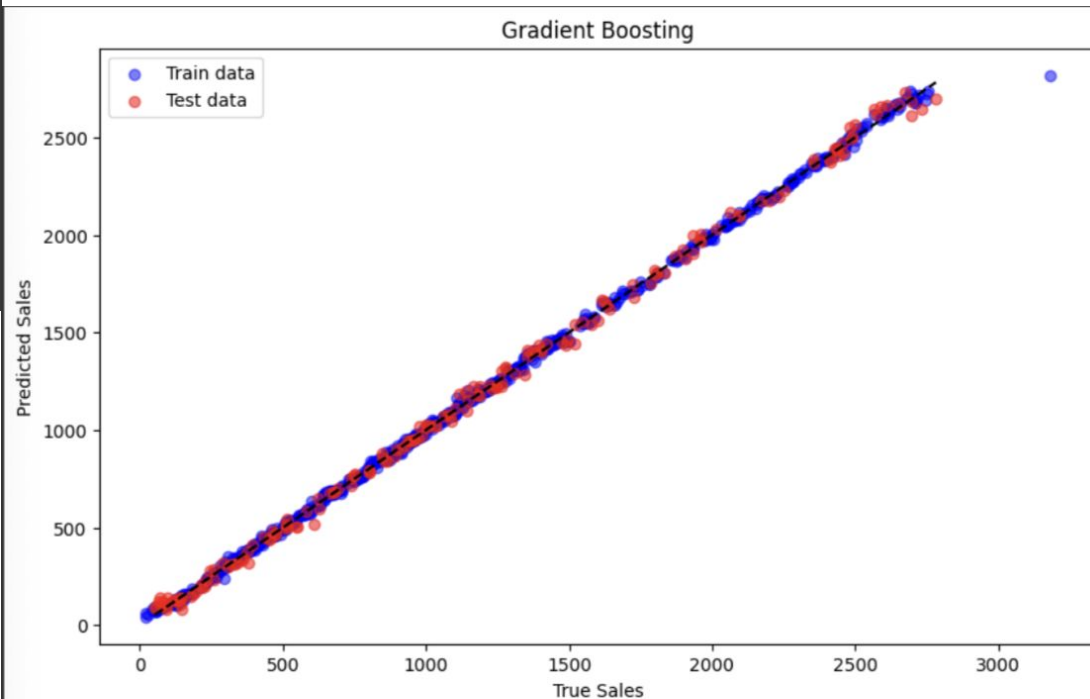
**Time Series Forecasting Models:** ARIMA, LSTM, XGBoosting, and Prophet. Used to forecast estimated sales for future (2025–2026). ARIMA (Autoregressive Integrated Moving Average) uses past values/errors to predict future values. LSTM (Long Short Term Memory) captured patterns in time series data and its a Recurrent Neural Network. XGBoost (Extreme Gradient Boosting) creates lag feature, rolling statistic and time related features. Prophet was developed for time series forecasting by Facebook.

Linear Regression Training MSE: 1.2154432126440646e-24, Testing MSE: 1.1103425  
Linear Regression Training R<sup>2</sup>: 1.0, Testing R<sup>2</sup>: 1.0



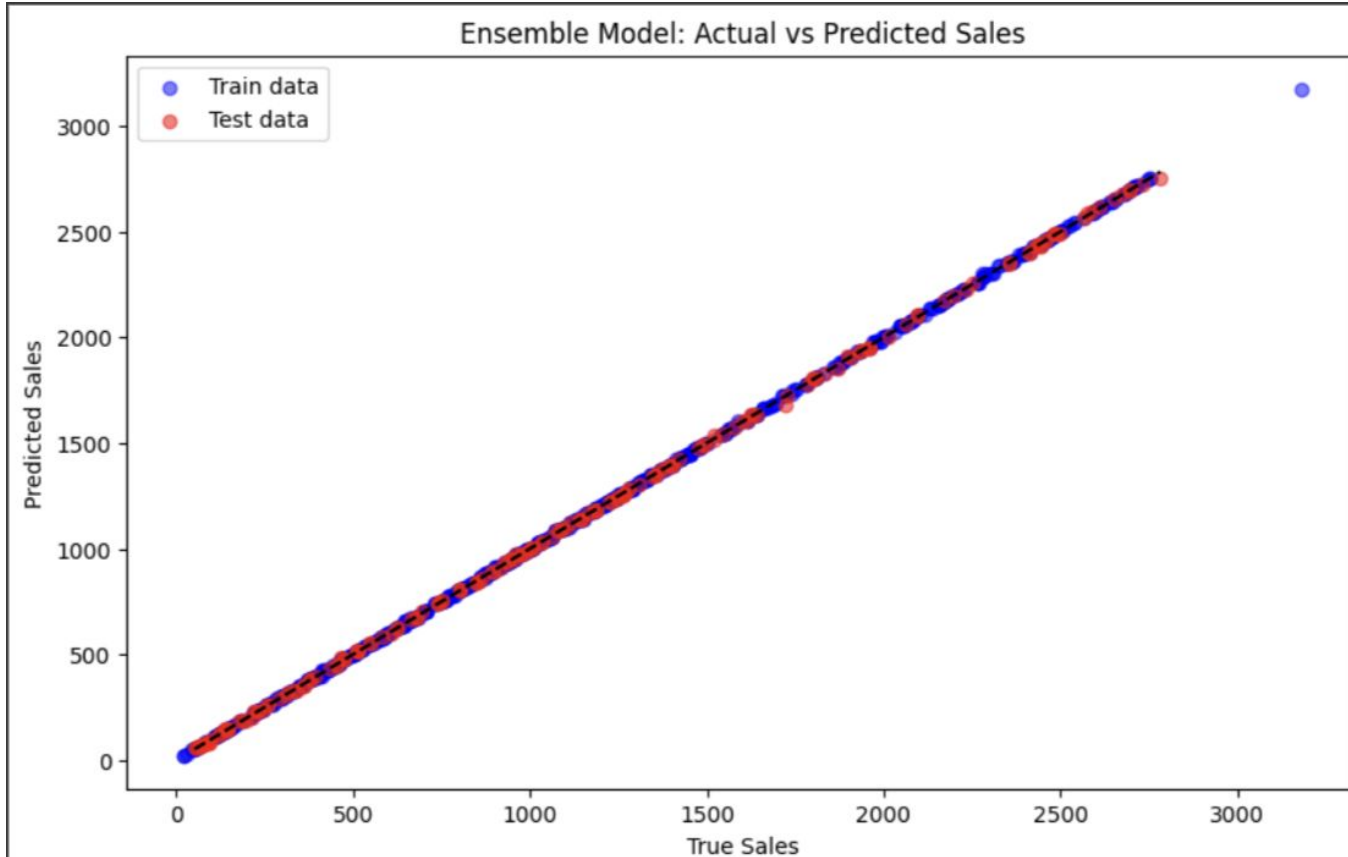
Random Forest Training MSE: 116.79235285359741, R<sup>2</sup>: 0.9997785220303747  
Random Forest Testing MSE: 422.1751271508969, R<sup>2</sup>: 0.9993103597859397  
Random Forest Cross-validated MSE: 968.32 (+/- 986.96)

# Demo of Models



Gradient Boosting Training MSE: 423.11296808273687, R<sup>2</sup>: 0.999197634101861  
Gradient Boosting Testing MSE: 1099.1822946247576, R<sup>2</sup>: 0.998204441085689  
Gradient Boosting Cross-validated MSE: 1500.59 (+/- 696.61)

# Demo of Models




Ensemble Model Training MSE: 20.984008550257567,  $R^2$ : 0.9999602071925536


Ensemble Model Testing MSE: 54.31190110669228,  $R^2$ : 0.9999112793040225

# (Demo) Forecasting with Regression Ensemble Model

Actual vs Predicted Sales (2019-2026)

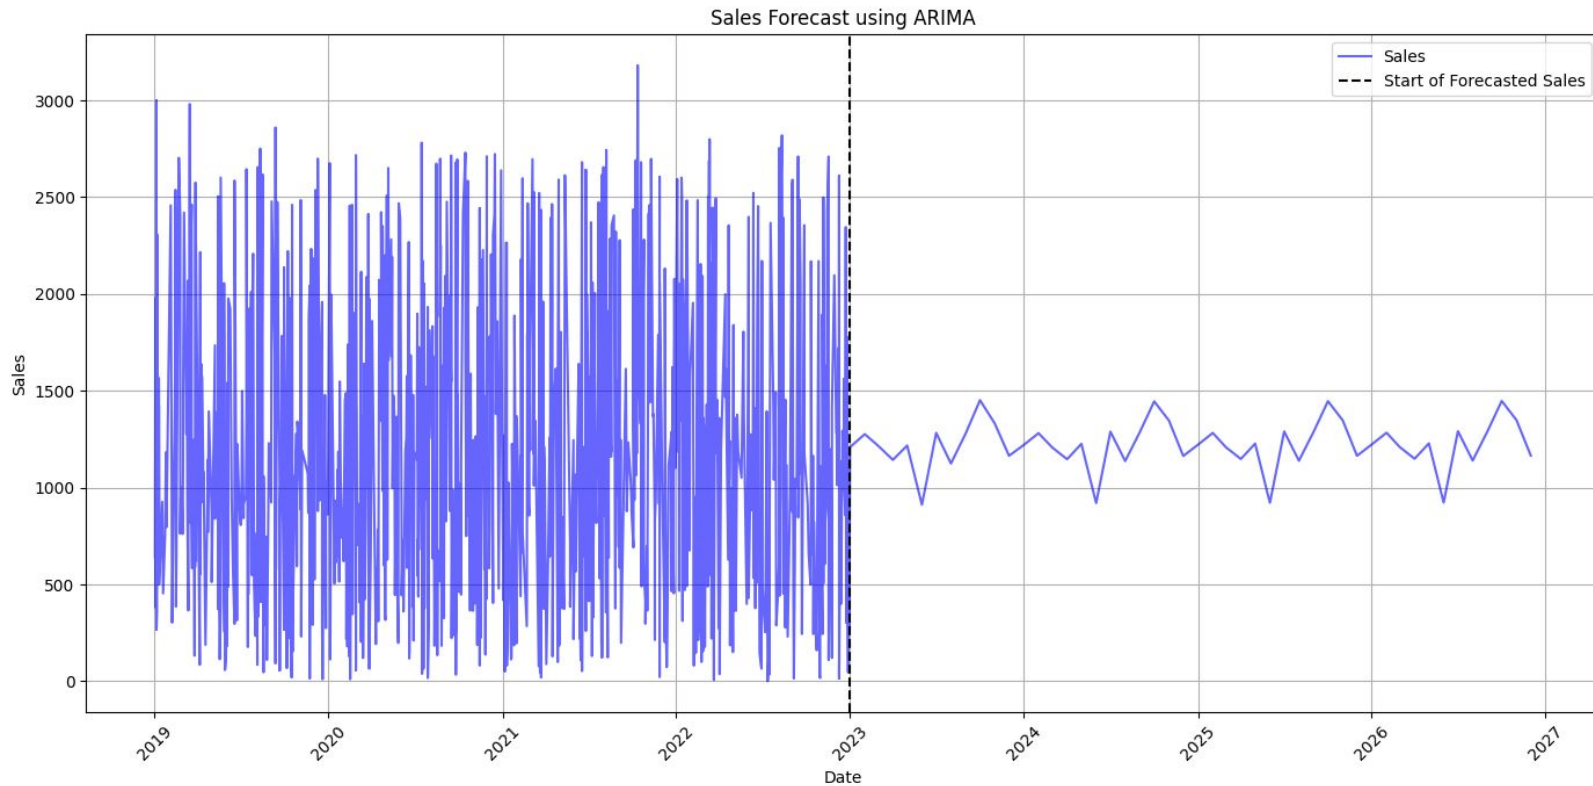


# (Demo) Forecasting with Time Series Model (Best Model)

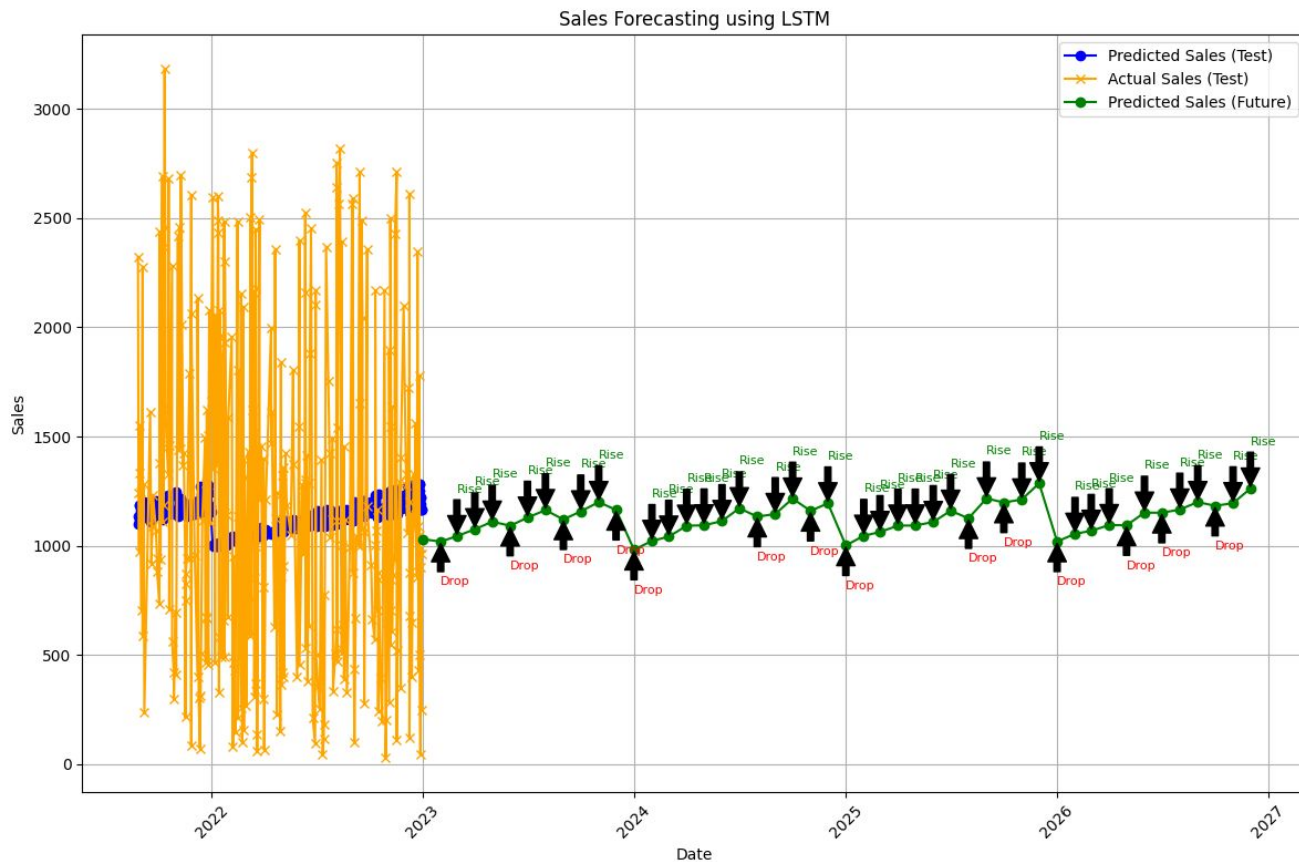




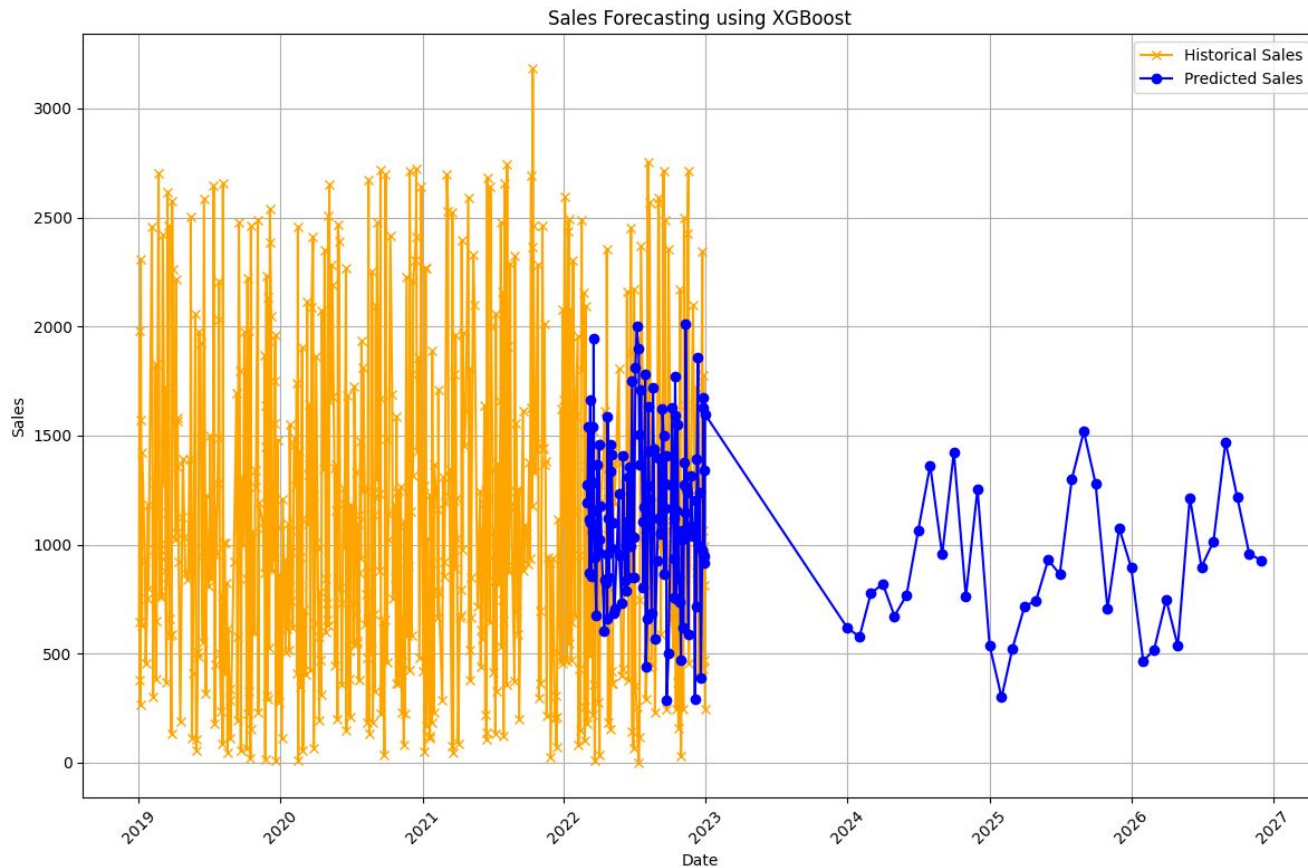
# (Demo) Forecasting with Time Series Model



# (Demo) Forecasting with Time Series Model



# (Demo) Forecasting with Time Series Model



# Models (Code Snippets)

```
34 # fit the Prophet model with holidays and additional regressors
35 model = Prophet(holidays=holidays, yearly_seasonality=True, weekly
36 # Add custom seasonalities for more depth and interperations
37 model.add_seasonality(name='monthly', period=30.5, fourier_order=5
38
39 # Add regressors (if any)
40 # model.add_regressor('additional_feature')
41
42 p_model = model.fit(train_data)
43
44 # Make predictions on the test set
45 future = model.make_future_dataframe(periods=len(test_data), freq=
46 forecast = model.predict(future)
47
48 # Plot the results
49 fig = plt.figure(figsize=(12, 8))
50 ax = fig.add_subplot(111)
51 fig = model.plot(forecast, ax=ax)
52 add_changepoints_to_plot(fig.gca(), model, forecast)
53 plt.plot(dataset['ds'], dataset['y'], 'o', markersize=2, label='Hi
54 plt.legend()
55 plt.xlabel('Date')
56 plt.ylabel('Sales')
57 plt.title('Sales Forecasting using Prophet')
58 plt.xticks(rotation=45)
59 plt.grid(True)
60 plt.tight_layout()
61 plt.show()
62
63 # Generate future dates for 2024, 2025, and 2026
64 future_dates = model.make_future_dataframe(periods=3*12, freq='M')
65 future_forecast = model.predict(future_dates)
```

```
9 def gradient_boosting(X_train, y_train, X_test, y_test):
10
11
12 # Convert datetime columns to numeric (Unix timestamp) before scaling
13 # for col in X_train.select_dtypes(include=['datetime64']).columns:
14 #     X_train[col] = X_train[col].astype(np.int64) // 10**9
15 #     X_test[col] = X_test[col].astype(np.int64) // 10**9
16
17
18 pipeline = Pipeline([
19     ('scaler', StandardScaler()), # Placeholder scaler, will be set in GridSearch
20     ('poly', PolynomialFeatures()), # Add polynomial features
21     ('regressor', GradientBoostingRegressor(random_state=42))
22 ])
23
24 param_grid = {
25     'scaler': [StandardScaler(), RobustScaler(), MinMaxScaler(), MaxAbsScaler()],
26     'poly_degree': [1], # Polynomial degree
27     'regressor_n_estimators': [100, 200, 300, 400, 500, 600, 700, 800, 900, 1000, 2000, 3000],
28     'regressor_learning_rate': [0.000001, 0.0000000000000001, 0.001, 0.0001, 0.01, 0.05,
29     'regressor_max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 20, 30, 40, 60, 70, 80, 90, 100], # M
30     'regressor_min_samples_split': [2, 5, 7, 9, 10, 12, 14, 20, 30, 40, 45, 50, 55, 60, 65, 70],
31     'regressor_min_samples_leaf': [1, 2, 8, 9, 10, 12, 15, 20, 30, 40, 50, 55, 65, 70, 75, 80, 90],
32     'regressor_subsample': [0.0001, 0.1, 0.2, .3, .4, .5, .6, 0.7, 0.8, 0.9, 1.0], # F
33 }
34
35 # Hyperparameter tuning with GridSearchCV
36 # search = GridSearchCV(pipeline, param_grid, cv=5, scoring='neg_mean_squared_error')
37
38 # search = RandomizedSearchCV(pipeline, param_grid, n_iter=50, cv=5, scoring='neg_
39 # search.fit(X_train, y_train)
40
41 # # Get the best model from the hyperparameter search
42 # best_model = search.best_estimator_
43 # print(f"Best parameters for Gradient Boosting Regressor: {search.best_params}")
44
45 # Perform cross-validation and get the mean score
46 # cv_scores = cross_val_score(best_model, X_train, y_train, cv=5, scoring='neg_mea
47
48 search = RandomizedSearchCV(pipeline, param_grid, cv=5, scoring='neg_mean_squared_
49 search.fit(X_train, y_train)
50 best_model = search.best_estimator_
51
52 # Train the best model on the entire training set
53 best_model.fit(X_train, y_train)
54
55 # Make predictions on the training and testing sets
56 train_predictions = best_model.predict(X_train)
57 test_predictions = best_model.predict(X_test)
```

```
17 # Ensemble function
18 def ensemble_model30(X_train, y_train, X_test, y_test):
19
20 """
21 Function to create and evaluate an ensemble model using stacking of multiple regression m
22 """
23
24 # Convert datetime columns to numeric (Unix timestamp)
25 for col in X_train.select_dtypes(include=['datetime64']).columns:
26     X_train[col] = X_train[col].astype(np.int64) // 10**9
27     X_test[col] = X_test[col].astype(np.int64) // 10**9
28
29 # Define base models
30 base_models = [
31     ('random_forest', random_forest(X_train, y_train, X_test, y_test)[0]),
32     ('gradient_boosting', gradient_boosting(X_train, y_train, X_test, y_test)[0]),
33     ('decision_tree', decision_tree(X_train, y_train, X_test, y_test)[0]),
34     ('neural_network', neural_network(X_train, y_train, X_test, y_test)[0]),
35     ('linear_regression', linear_regression_model(X_train, y_train, X_test, y_test)[0]),
36     ('elastic_net', elastic_net(X_train, y_train, X_test, y_test)[0]),
37     ('lasso_regression', lasso_regression_model(X_train, y_train, X_test, y_test)[0]),
38     ('ridge_regression', ridge_regression(X_train, y_train, X_test, y_test)[0])
39 ]
40
41 # Define stacking regressor
42 ensemble = StackingRegressor(
43     estimators=base_models,
44     final_estimator=GradientBoostingRegressor(), # Final estimator
45     passthrough=True
46 )
47
48 # Train the ensemble model
49 logging.info("Training the ensemble model...")
50 ensemble.fit(X_train, y_train)
51
52 # Make predictions
53 train_predictions = ensemble.predict(X_train)
54 test_predictions = ensemble.predict(X_test)
55
```

# Conclusion

## Results and Limitations:

- The accuracy of the forecasting of Walmart sales for the future depends on many factors. Highest estimated sales were in the range of 0-3000\$, most products were home appliances excluding other products sold at Walmart. Used 1000 samples due to execution time for models. Weren't able to find exact real trends specifically for US, graphs, and other specifics to compare with our findings.
- Researched and coded using various websites such as stackoverflow, geeksforgeeks, and others however we did struggle to find specific tutorials on complex coding areas. Used Google Colab but were limited on which GPU to use.

## Key Findings:

- Linear Regression, Ensemble, and Prophet models, were the only three that gave us the best in fitting on data and predicting. The two regression models had better  $R^2$  and MSE. These two models can do good on unseen data. We were surprised LR did a bit better than Random Forest and Gradient Boosting, we were expecting RF and GB to do better after researching about it on complex data.
- Prophet was the only model out of the other 3 time series model that gave us a better and accurate results. We spent a lot of time trying to work with Arima, LSTM, XGBoosting but these three gave us more confusing forecasting which look off from the historical sales. But Prophet gave us a easier interpretations and better representation.
- Ensemble and Prophet model do have some similarities only in the trends of sales.



# Conclusion

Model	Training MSE	Testing MSE	Training R <sup>2</sup>	Testing R <sup>2</sup>
Lasso Regression	0.4399621278	3.3189848872	0.9999992245	0.9999933602
Random Forest	143.9248826342	4838.8292807894	0.9997463171	0.9903197451
Gradient Boosting Model	5574.9082003059	4134.8262450875	0.9901736321	0.9917281289
Decision Tree	0.8602539747	3586.6020339802	0.9999984837	0.9928248715
Neural Network	951.6560332385	11892.8469081357	0.9981953369	0.9805725516
Linear Regression	1.2154432126	1.110342550026471	1.0000000000	1.0000000000
Elastic Net	4.6861899849	2.6315806283	0.9999911134	0.9999957012
Ridge Regression	1.1381866	1.388515542	1.0000000000	1.0000000000
Ensemble Model	20.9840085503	54.3119011067	0.9999602072	0.9999112793

Model	Training MSE	Testing MSE	Training R <sup>2</sup>	Testing R <sup>2</sup>
Lasso Regression	0.4399621278	3.3189848872	0.9999992245	0.9999933602
Random Forest	143.9248826342	4838.8292807894	0.9997463171	0.9903197451
Gradient Boosting Model	5574.9082003059	4134.8262450875	0.9901736321	0.9917281289
Decision Tree	0.8602539747	3586.6020339802	0.9999984837	0.9928248715
Neural Network	2345.0347563957	10289.5367098171	0.9958666272	0.9794154056
Linear Regression	5.951888742	7.3544236781	1.0000000000	1.0000000000
Elastic Net	155.8785344817	203.1706101532	0.9997252475	0.9995935498
Ridge Regression	1.17254906	2.389266	1.0000000000	1.0000000000
Ensemble Model	20.5206590949	343.5874968719	0.9999638302	0.9993126406

# Conclusion

## Mistakes:

- We made mistakes during this project. Spent too much finding the right dataset, didn't research enough on specific areas until later, experimented too much, changed our goal many times due to limited dataset causing a redo on the project, confused on certain steps, specified or generalized the goal, didn't use the best IDE, and took time to realize how to decrease execution times so we could adjust the code for better results.

## What we learned:

- We learned to first have a better plan or outline of the project, to stick with it, use better IDE, and do more research before coding. To write code in blocks step by step, better for readability, and commenting. Learned alot about the models, how to use these tools from the libraries, and how to research. Understood concepts of machine learning.

## What's for the future?:

- For future work we plan to make our goals more complex in this area of predictions and analysis of data and working with it. To focus more on these models, explore advanced models, and expand on this project to make improvements for career learning.



# Q&A

Thank you for your attention !!!

Why did the predictive model apply for a job? It wanted to improve its training.

And why was the predictive model always relaxed? Because it never overfitted.

But why was the predictive model always invited to the strategy meetings? Because it could see the future!

## Questions?

